JURIAID – AN AI-POWERED LEGAL RESEARCH AND EVIDENCE ASSISTANT

Project ID : 25-26J-274

Project Proposal Report

Prarthana A.P.S – IT22127228

Supervised by: Mr. Jeewaka Perera

Co-supervised by: Ms. Madusha Weerasooriya

External Supervised by: Mrs. Nadeeka Senavirathne

BSc. (Hons) Degree in Information Technology Specializing in Information Technology

Department of Information Technology

Sri Lanka Institute of Information Technology

Sri Lanka

July 2025

JURIAID – AN AI-POWERED LEGAL RESEARCH AND EVIDENCE ASSISTANT

Project ID : 25-26J-274

Project Proposal Report

BSc. (Hons) Degree in Information Technology Specializing in Information Technology

Department of Information Technology

Sri Lanka Institute of Information Technology

Sri Lanka

July 2025

DECLARATION

We are confirming that the work given in this proposal is one of our own volition. The same work has not been given, entirely or partially, for academic credit at any other university or institution. We have also made sure that all sources of information and ideas from others have been clearly cited and appropriately acknowledged in the document.

Name	Student ID	Signature
Prarthana A.P.S	IT22127228	

The above candidates are researching the undergraduate Dissertation under my supervision.

Name of supervisor: Mr. Jeewaka Perera

Name of co-supervisor: Ms. Madusha Weerasooriya

Signature of supervisor: Date: 2025/08/29

Signature of co-supervisor: Date: 2025/08/29

ABSTRACT

This project removes a serious inefficiency of Sri Lankan legal practice: the time-consuming, manual process of looking up, verifying, and placing in context relevant statutes and amendments to them. Present legal databases offer only skeletal keyword search, have no semantic understanding, and perform no automatic follow-through on law changes, and therefore force lawyers to devote 6–10 hours per case to statutory research. Our proposed solution is LawStatKG, which is an AI-based legal retrieval system that streamlines and automates statute identification. It extracts relevant Sri Lankan statutes and amendments through hybrid retrieval (BM25 + LegalBERT embeddings), a legal knowledge graph (RDF/Neo4j), and temporal tracking of amendments. It accepts natural language English queries, retrieves conceptually related law, and brings out automatically all relevant amendments and cross-references. The core design employs LegalBERT optimized for semantic understanding, a legal domain-specific RDF ontology to model legal relations, and a version-aware timeline system to track the provenance of laws. This enables the system to offer courtroom-ready, rolled-up output with the statute text, amendment history, and referenced provisions in one view. By automating the essential legal task of fetching statutes, the research aims to reduce case preparation time significantly, decrease legal errors, and make access to accurate statutory information democratised. The system emphasizes reliability, usability, and interworking with the common AI legal assistant platform to render it realistically deployable in the Sri Lankan legal environment.

Keywords: Legal AI, Statute Retrieval, Hybrid Search, Legal Knowledge Graph, Temporal Tracking, Legal-BERT, Legal Tech, Sri Lanka

TABLE OF CONTENTS

DECLARATION 3
ABSTRACT 4
TABLE OF CONTENTS 5
LIST OF FIGURES 6
LIST OF TABLES 7
LIST OF ABBREVIATIONS 8
1. INTRODUCTION 9
1.1 BACKGROUND 9
1.2 LITERATURE SURVEY 10
1.3 RESEARCH GAP 12
1.4 RESEARCH PROBLEM 14
2. RESEARCH OBJECTIVES 15
2.1 MAIN OBJECTIVE 15
2.2 SPECIFIC OBJECTIVES 15
3. METHODOLOGY 16
3.5 Project Flow 19
3.6 Commercialization 20
4. PROJECT REQUIREMENTS 21
4.1 Functional Requirements 21
4.2 Non-functional Requirements 21
4.3 User Requirements 21
5. BUDGET DETAILS 22
REFERENCES 23
APPENDICES 24
Appendix 1. Work Breakdown Structure (WBS) 24
Appendix 2. Gantt Chart 25
Appendix 3. Similarity Report 26

LIST OF FIGURES

Figure 2.3.1 System Overview Diagram..... 17

Figure 3.3.2 Component system diagram 18

LIST OF TABLES

Table 1.3.1 Comparative analysis highlighting the novelty of the proposed system 12

Table 2.1 Budget for the proposed system 22

LIST OF ABBREVIATIONS

ABBREVIATION	DESCRIPTION
LawStatKG	Law Statute Knowledge Graph
KG	Knowledge Graph
LLM	Large Language Model
BM25	Best Match 25
CTV	Component Temporal Version
TV	Temporal Version
TF-IDF	Term Frequency-Inverse Document Frequency
NLP	Natural Language Processing
e-Laws	Electronic Laws Portal
SPARQL	SPARQL Protocol and RDF Query Language
NLL	Natural Language Leagal

1. INTRODUCTION

1.1 BACKGROUND

Punctuality, accuracy, and scale of statutory research in modern legal practice directly impact the quality of legal representation and administration of justice. Legal researchers, lawyers, and judges entirely depend on timely and accurate availability of relevant laws, amendments, and associated statutes during preparation of legal arguments or judicial decisions. Delayedness or inaccuracy in identification of provisions of law that are applicable can lead to weakening of legal submissions, procedural delay, or denial of justice. Statutory research in Sri Lanka is primarily, manual, and laborious. The process typically consists of initial search for the main legislation that regulates a case, and subsequent comprehensive verifications of its present validity, applicable amendments, repeals, and cross-references to related acts. Lawyers usually rely mainly on the e-Laws portal, gazettes, and bulk PDF repositories. These machines tend to operate almost completely on word-level searches, which tend not to capture the underlying legal intent or word-level nuances. Second, amendments also have to be complied with manually by cross-verifying different sources, and there is no facility of an automatic system encompassing all the content incorporating its reasonable integration. It is a labor-intensive process, prone to leaving something out, and particularly inconvenient for rural practitioners who have no easy access to up-to-date legal materials and networks of expertise. The absence of semantic understanding, relationship mapping, and history tracking in existing legal research tools is thus evidently a technology deficit.

This is especially difficult in a common law jurisdiction like Sri Lanka's, where statutory interpretation requires not only consideration of the content of the statute itself but also of its history, context, and relationship to other statutes. In reality, this means lawyers taking between six to ten hours to prepare for a single case, with corresponding high probabilities of failing to catch relevant amendments or interfacing provisions, particularly for complex legal issues. This research proposes LawStatKG, a legal statute knowledge graph module with AI integrated into a larger legal assistant software.

Distinct from conventional search engines, LawStatKG combines semantic retrieval, relation-aware navigation, and temporal versioning to provide an ever-changing map of Sri Lankan law. It will recognize natural language questions in law, import relevant statutes using a hybrid approach combining lexical search through BM25 with searching by semantic embeddings using LegalBERT, and link each statute to its amendments and related acts into a knowledge base specifically adapted to RDF. The system will also track historical development of provisions in law, showing how laws have developed over time and when each amendment became active, and constructing aggregated, attorney-level reports that combine contemporary statutory language, pertinent amendments, and coterminous laws into one last authoritative reference. By transforming isolated legal knowledge into a smart, consolidated, and searchable

set of information, LawStatKG will cut research time by a significant percentage, equalize access to good quality legal information, and make case preparation more accurate and more so. It will bridge the urban–rural divide in legal research capability and help build Sri Lanka's current legal information infrastructure. Lastly, the aim of this study is to show how these synergies between semantic search, knowledge graph technologies, and natural language processing can make individuals more accessible to justice, enhance the practice of law, and pay long-term dividends to the legal community.

1.2 LITERATURE SURVEY

The use of Artificial Intelligence in legal information systems has progressed significantly with the evolution of Retrieval-Augmented Generation (RAG), a system that harmonizes external knowledge retrieval with the generative capabilities of Large Language Models (LLMs). This is consistent with how it addresses the inherent vulnerabilities of standalone LLMs—such as hallucinations, out-of-date information, and untraceable reasoning—by anchoring answers to verifiable, external proof, hence making it highly suitable for high-risk applications like law [1].

Hybrid Retrieval: Merging Different Search Paradigms for Robustness

Yet another crucial advancement in RAG systems is the implementation of hybrid retrieval, when different paradigms of retrieval are brought together to offer more robustness and accuracy. While Gao et al. present a comprehensive review of RAG advancements, they note that naive retrieval (BM25 or dense embeddings) generally fails to achieve a balance of precision and recall in complex domains [1]. To oppose this, Kalra et al. propose HyPA-RAG, a system that integrates sparse (BM25), dense (embedding-based), and knowledge graph (KG)-based retrieval into a single pipeline [2]. This combination allows the system to utilize keyword matching, semantic semantics, and relational reasoning simultaneously. The effectiveness of hybrid retrieval is also confirmed by Sheetal et al., who experiment with TF-IDF, BM25, and Sentence-BERT for legal document similarity, noticing that while each method has its strengths, their combination yields more consistent results [3]. Similarly, Silva et al. demonstrate in Graph RAG how the retrieval performance is boosted when the structural as well as semantic information are merged, although they don't use BM25 [4]. These findings collectively suggest that there is no one retrieval strategy that will effectively work for legal AI, and that a hybrid multi-strategy approach is required to obtain reliable results.

Knowledge Graphs: Enabling Relational and Contextual Reasoning

Knowledge Graphs (KGs) in RAG systems have now become the first-order innovation for contextual intelligence improvement. Kalra et al. provide a KG of policy objects (requirements, obligations) and introduce KG triplets as additional context to the LLM, improving traceability

and explainability of outputs [2]. This allows the system to reason about legal concept relationships instead of isolated text.

Sheetal et al. follow the same path by developing a Neo4j-based KG of Indian court judgments, in which rhetorical roles (Facts, Ruling, Precedent) are nodes and relation between them is edges [3]. Their study demonstrates that the graph structure of legal text allows for thematic similarity computation and makes case law recommendation systems possible. Silva et al. extend this idea in Graph RAG, where legal knowledge graph is not only a fixed structure but a dynamic, hierarchical graph that includes legal norms, components, and versions [4]. It allows for structure-aware retrieval where queries can be targeted to certain levels (Article, Section) of a statute. Collectively, these works demonstrate that KGs are not repository software but dynamic entities that augment retrieval, reasoning, and response generation in legal AI.

Temporal and Structural Modeling: Modeling Legal Evolution

Legal texts are dynamic; they evolve over time by repeals, amendments, and revisions. Silva et al. address this in Graph RAG by introducing Temporal Version (CTV) nodes, a snapshot of a legal provision at a particular point of time [4]. This enables point-in-time queries critical property for correct legal interpretation. The model also handles alteration actions ("amended by Act Z") as first-class things, so the system has a record of the history in which the laws evolved. Even though Gao et al. and Kalra et al. do not model temporal change directly, their designs do support versioned data integration in retrieval and augmentation processes [1], [2]. Neither Sheetal et al. supports temporal tracking because they're more concerned with judgment similarity than statutory change [3]. The absence of temporal modeling in most systems, however, highlights the innovation of Graph RAG in addressing the dynamic nature of law.

Though these four papers are a huge leap forward—Gao et al. [1] extending RAG, Kalra et al. [2] suggesting hybrid retrieval, Sheetal et al. [3] employing KGs to model legal judgments, and Silva et al. [4] advancing structural and temporal modeling—none of them are for Sri Lankan law. There is no system that combines hybrid retrieval, a Sri Lankan statute legal knowledge graph, automated amendment linking, and temporal changes tracking into an end-to-end pipeline yet.

This work fills that gap by presenting LawStatKG, the first Sri Lanka-centered legal retrieval engine that benefits from the strengths of these approaches. By hybrid retrieval over HyPA-RAG, KG methods of Indian legal NLP, and temporal-structural graph modeling from Graph RAG, the system presented here offers a valid, context-tailored, and timely answer for statutory research in Sri Lanka.

1.3 RESEARCH GAP

Despite the incredible advancements in Retrieval-Augmented Generation (RAG) and its implementation in legal applications, there remains one overwhelming research gap: the absence of a hybrid, knowledge-driven, time-sensitive legal retrieval system suitable for the Sri Lankan statutory framework. While the four said works—Gao et al. [1], Kalra et al. [2], Sheetal et al. [3], and Silva et al. [4] - collectively represent the state of the art in RAG for legal purposes, none of them takes into account the pertinent structural, linguistic, and procedural characteristics of Sri Lankan law, and there is no existing system that encompasses all building blocks needed to facilitate proper, context-dependent, and historically informed legal research in the jurisdiction. A look at the literature reveals that key capabilities have been developed in isolation but not yet integrated for Sri Lanka.

To identify the research gap, we compared four research papers in more recent works to our solution. The results are shown in Table 1.3.1.

Feature	Research (Y.Gao) [1]	Research (R.Kalra) [2]	Research (S.sheetal) [3]	Research (J.silver) [4]	Proposed System
Hybrid Retrieval (Combines TF-IDF keyword search and embeddings)	Discussed. But not implement	Yes	No	Partial	Yes
Legal Knowledge Graph	No	Yes	Yes	Yes	Yes
Automated Amendment Linking	No	No	No	Yes	Yes
Temporal Tracking of Legal Changes	No	No	No	Yes	Yes
Jurisdiction-Specific Design	No	Yes- USA	Yes- India	Yes - Brazile	Yes -Sri lanka
Combine System	No	No	No	Yes	Yes

Table 1.3.1 Comparative analysis highlighting the novelty of the proposed system

1. Hybrid Retrieval is Proven, but Not Yet for Sri Lankan Law

The hybrid retrieval concept—combining sparse (e.g., BM25), dense (e.g., embeddings), and knowledge-based methods—is proven. Gao et al. [1] summarize RAG's development and assert the superiority of hybrid over one-method retrieval. Kalra et al. [2] demonstrate this in practice with HyPA-RAG, which integrates BM25, semantic embeddings, and knowledge graph triplets to facilitate higher accuracy and prevention of hallucinations in legal systems for AI. However, their system is designed for NYC Local Law 144, not Sri Lankan legislation. While the method is established, it has not been applied to the Commonwealth-form legislative code, multilingual setting (English, Sinhala, Tamil), or Sri Lankan juridical vocabulary.

2. Knowledge Graphs Add Context, but Aim Case Law, Not Statutes

Using knowledge graphs (KGs) to harvest legal relations is a familiar refrain. Sheetal et al. [3] construct a Neo4j-based knowledge graph (KG) of Indian Supreme Court judgments, where rhetorical roles (e.g., Facts, Ruling, Precedent) are nodes and thematic similarity between nodes are edges. This enables more intelligent case recommendations and similar computation. Kalra et al. [2] also uses a KG to identify relevant policy entities and inter-entities. Both systems are interested in judicial rulings or policy compliance, not statutory law. There is no KG that represents Sri Lankan Acts, Sections, Amendments, and their interconnection - a primary need for statutory research.

3. Temporal and Structural Modeling is Advanced but Geographically Limited

The most impactful structural and temporal modeling innovation is attributed to Silva et al. [4], who propose Graph RAG. This method depicts legal norms as a hierarchical and temporal graph with Component Temporal Version (CTV) nodes containing the text of a provision at a time. This facilitates point-in-time queries - a fundamental requirement for determining the law on a previous date. Their contribution is novel, albeit experimented upon the Brazilian legislation system, which has a unique legislative drafting style from the one in Sri Lanka based on the Commonwealth. The technique is sound, but not in the local mode.

4. No Integrated Sri Lankan Statutory Research System

Current tools made available to Sri Lankan legal practitioners, e.g., the e-Laws portal, offer only rudimentary keyword searching without:

Semantic understanding of legal concepts,

Cross-referencing of amendments done automatically,

Relationship mapping between enactments,

Or historical version tracking.

1.4 RESEARCH PROBLEM

The main problem this research attempts to address is that there does not exist any reliable way for Sri Lankan lawyers and legal researchers to quickly and accurately retrieve the current or past version of the law. Sri Lankan lawyers, judges, and legal researchers now must check manually at the time of the current moment and past version of laws. They rely on the e-Laws portal, which supports only rudimentary keyword searching and lacks automated provisions to monitor amendments, merge updated provisions or find the very specific wording of a law as of a specific date.

While the current hybrid retrieval systems have made the integration of semantic search, knowledge graphs, and temporal tracking more likely in legal analysis, such systems are available in the foreign jurisdictions (USA, Brazil) and are not available or implemented for Sri Lanka law. This paper aims to overcome this limitation by developing LawStatKG, the first legal retrieval engine specifically designed for Sri Lanka that combines:

1. Hybrid retrieval (BM25 + LegalBERT) for precise search.
2. A legal knowledge graph to map the relationships among amendments and laws,
3. Temporal versioning to follow the evolution of the laws over time, And automated combination of current and past law books.

RESEARCH QUESTION

There does not exist any reliable way for Sri Lankan lawyers and legal researchers to quickly and accurately retrieve the current or past version of the law.

2. RESEARCH OBJECTIVES

2.1 MAIN OBJECTIVE

To create and implement LawStatKG, a hybrid, knowledge-based, and temporal legal search environment for Sri Lankan statute law. The environment will include semantic search, legal knowledge graph, and temporal versioning to allow practitioners to automatically find up-to-date and historical laws, remove manual labor, remove errors, and improve access to accurate legal information throughout Sri Lanka.

2.2 SPECIFIC OBJECTIVES

1. **To develop a hybrid retrieval engine that incorporates keyword-based, semantic, and knowledge graph-based methods for precise legal search.** This is achieved by constructing a retrieval pipeline that integrates BM25 for robust keyword match, LegalBERT embeddings for understanding legal term meaning, and a knowledge graph (KG) retriever to extract suitable legal relationships. The results of the three methods will be combined by Reciprocal Rank Fusion (RRF) to produce a better and more comprehensive result set than any single approach alone.

2. **To construct a Sri Lanka-specific legal knowledge graph (LKG) that would replicate the hierarchical structure of statutes and their interrelations.** This involves building an RDF/Neo4j-powered graph in which legal entities such as Acts, Parts, Chapters, and Sections are considered nodes, and relationships between them (amends, repeals, cites) as edges. The graph will enable "semantic navigation" to allow users to travel from a general subject to a specific clause in legislation, such that the context returned is comprehensive according to the law.

3. **To enable a temporal versioning system capturing the evolution history of modifications in the legislation over time.**
   That is, every legal provision will be represented using Temporal Version (CTV) nodes, which contain a precise copy of a law at a specific date. This framework will automatically propagate, so that whenever an update is done to a section, all its parent pieces (articles, chapters) that include it get updated to a fresh version. This is done by generating Text Units for each legal provision, including not just the isolated text of an individual section, but also the necessary context of its parent pieces (the capture of an article). The Text Units will be used for retrieval and provided to the LLM for the

purpose of generating a final output with the statute text, its amendment history, and surrounding provisions, in a human-readable, easy-to-understand format.

3. METHODOLOGY

Development of LawStatKG will be performed employing a systematic approach with five broad phases: Data Acquisition, preprocessing, model construction, knowledge graph construction, and system integration and testing. The approach will combine Natural Language Processing (NLP), semantic search techniques, and knowledge representation models to deliver a time-critical, relationship-sensitive statute retrieval system tailored to Sri Lankan law.

OVERALL SYSTEM DIAGRAM

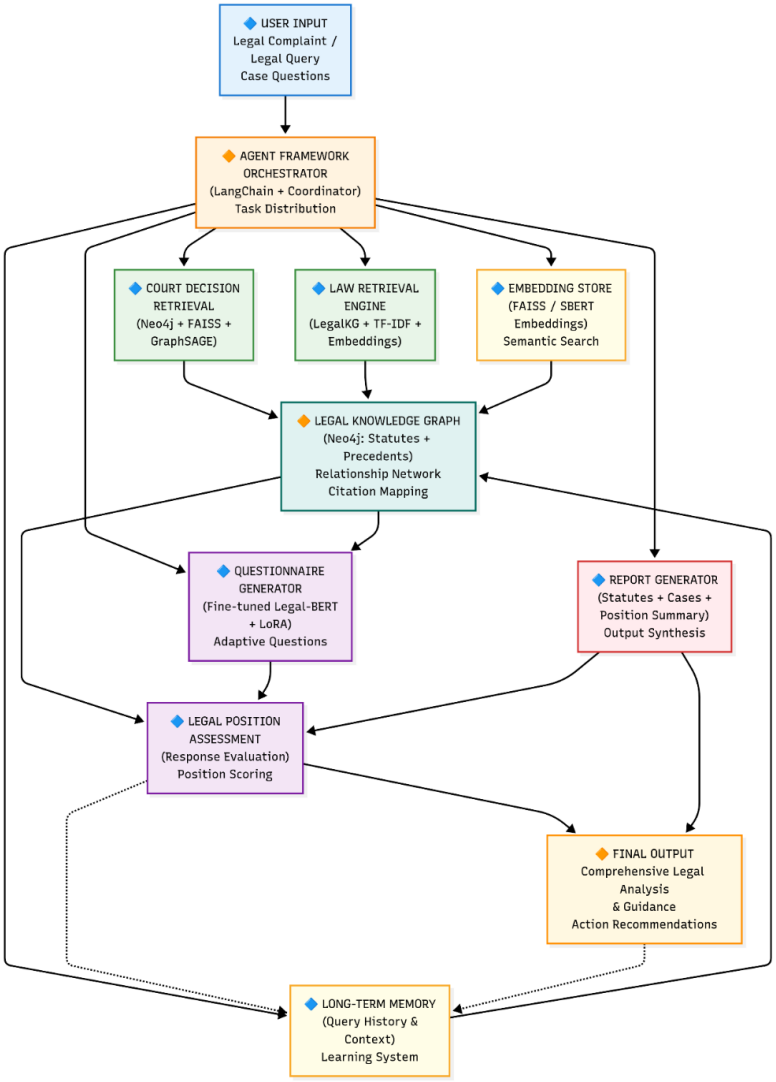

Figure 1.3.1 System Overview Diagram

MY COMPONENT SYSTEM DIAGRAM

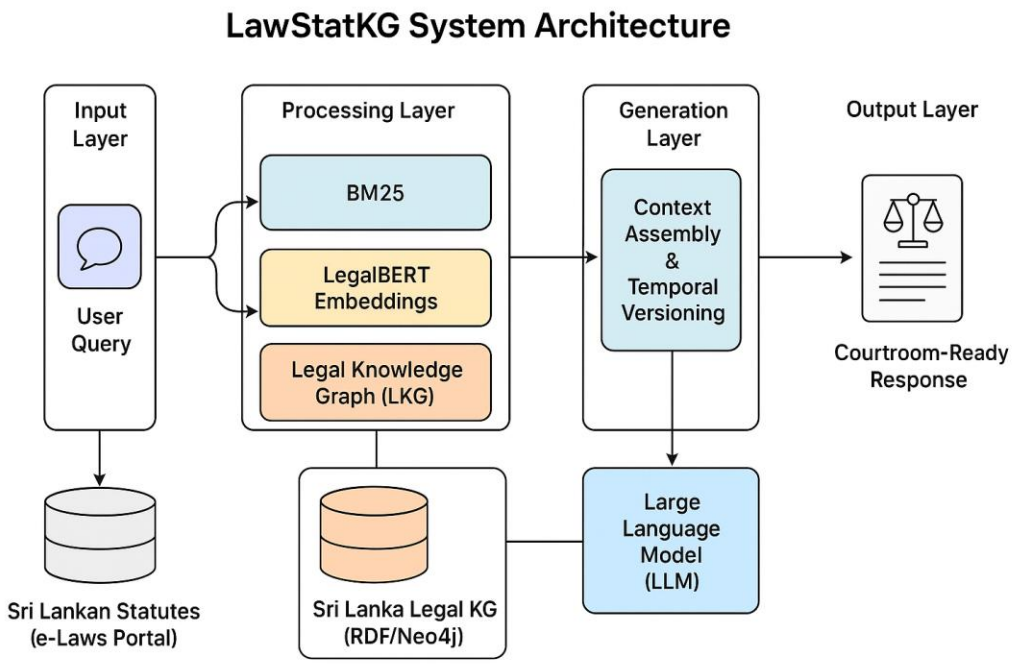

Figure 2.3.2 Component system diagram

LawStatKG — System Architecture Overview

1. Input Layer

NLL Query (laws that govern illegal land occupation)

Filter Options (jurisdiction, date range)

2. Data Sources Layer

Statute Corpus (text-based form of the laws)

Gazette Amendments Database

Related Acts Repository

Metadata Store (jurisdictions, dates, categories)

3. Processing Layer

Query Preprocessing

Hybrid Retrieval Engine

BM25 Lexical Search (match retrieval)

LegalBERT Semantic Search (meaning-based retrieval)

Knowledge Graph Access Layer

RDF/Neo4j Database Query

Entity & Relationship Retrieval

Temporal Tracking Module (history of versions, dates of enactment)

4. Output Layer

Consolidated Statute View (incorporated with amendments)

Relationship Map Visualization (graphical output of legislations with relationship)

Temporal Timeline View (historical evolution of the law)

Courtroom-Ready Report (PDF/HTML output).

3.5 Project Flow

1. Data Acquisition

Firstly, collect authoritative evidence of enactment, Major enactments of the e-Laws web site of the Sri Lanka Parliament. Gazette notices including amendment, repeal, and subsidiary legislation are noticed. Law reports and PDF archives available to the general public were connected to legal sources. PDF-to-text and web scraping will be employed if necessary to form a homogeneous corpus of law. All of the sources will be stored permanently in a metadata repository with publication date, legal topic, and gazette number so that temporality can be easily tracked when it must be done in the future.

2. Preprocessing and Normalization

Preprocessing cannot be avoided because of heterogeneity of Sri Lankan legal data. Steps are: Optical Character Recognition (OCR) of scanned gazette pages. Text normalization and cleaning (removing formatting artifacts, OCR correction, and terminological standardization.) Segmentation of the laws into clauses, subsections, and sections. Metadata tagging, enactment date, amendment date, and related laws, using rule-based and machine learning algorithms.

3. Development of Hybrid Retrieval Engine

The retrieval sub-system will have: Hybrid search engine: Lexical search through BM25 for exact keyword matches. Semantic search through a Sri Lankan legal text trained LegalBERT model for contextual sense. Query interpretation to execute natural language

queries and map them onto relevant legal provisions. Ranking fusion to combine lexical and semantic scores to get best retrieval precision.

4. Legal Knowledge Graph Construction

We will create a domain-specific RDF-based knowledge graph to model legal relations. Ontology building to define entities (statute, amendment, provision) and relations (has amendment, refers to, related to). Entity linking to link statutes to their amendments, related acts, and cross-references. Temporal versioning to store historic amendments, hence being capable of querying the laws on a specific date. Graph database deployment via Neo4j or an RDF with SPARQL query support.

5. Consolidation and Report Generation

The software must have an auto-document composition facility that: Combines the initial legislation with all the appropriate amendments. Places the legislation as it stands currently, its previous changes at will are available. Generates court-proof PDF or HTML output with references, cross-references, and context.

6. Testing

This process is iterative. Feedback obtained through the user testing in the last step may initiate changes in the earlier steps, i.e., adjusting the knowledge graph scheme or retraining LegalBERT model.

Testing comprises:

Unit Testing: Component-level testing

Integration Testing: Entire system is tested

Real-World Testing: Operating in law environment.

User Testing: Lawyers' usability and feedback

3.6 Commercialization

The advice would have been that such a system be launched as a stand-alone entity, independent of a cooperative organization, mainly because there is no centralized efficient platform in Sri Lanka that connects legal professionals to AI tools. Any lawyer, regardless of their firm or affiliation, would be free to use the system this way. Multi-device system design will be the next step with a mobile app for lawyers to access the system's features. Probably, this product will follow a freemium business model. The free version could have

the basic functionality of generating questions or set a limit on the number of questions per case. The premium version would come with greatly enhanced features such as:

• Unlimited and elaborate question generation that deepens the legal context and nuance.

• This connectivity feature expands the legal knowledge base along with real-time legal updates.

• The advanced fact elicitation and iterative refinement provide the user with additional assistance in research for the complimentary legal practice.

4. PROJECT REQUIREMENTS

4.1 Functional Requirements

1. Keyword (BM25) and semantic (LegalBERT) search of statutes.
2. Knowledge graph of statute relationship (amendments, references) to store.
3. Show historical versions of statutes to monitor.
4. Show relationship-aware results with graph visualization.
5. Generate one, courtroom-ready report with statute text, amendments, and relationships.
6. Export as PDF or Markdown.

4.2 Non-functional Requirements

1. Performance – Show results in less than 3 seconds.
2. Scalability – Handle larger statutes efficiently.
3. Security – Use web based and role-based access control.
4. Reliability – Give 99% uptime supported with backups.
5. Maintainability – Documentation and Modular development
6. Integration – Used REST APIs.
7. Concurrency – Handle 500+ concurrent users.

4.3 User Requirements

1. Reduce statute research time from hours to minute.
2. Provide up-to-date laws with clear amendment history.
3. Claims to comprehend legal principles, not keywords alone.
4. Demonstrates how laws relate to each other in an easy-to-understand visual way.
5. Generates professional, ready-to-use output for court.

5. BUDGET DETAILS

Description	Cost	Occurring
Server Hosting	Covered BY GitHub Pack	Monthly
Cloud Database	Covered BY GitHub Pack	Monthly
Legal Metirials	LKR 10,000	Annual
Developer Tools/Software Subscriptions	Covered BY GitHub Pack	Monthly/Annual
LLM API usage	LKR 10,000	Monthly

Table 2.1 Budget for the proposed system

REFERENCES

[1] Z. W. A. G. A. H. X. G. A. K. a. P. T. R. Kalra, "HyPA-RAG: A Hybrid Parameter Adaptive Retrieval-Augmented Generation System for AI Legal and Policy Applications," London, 2024.

[2] R. B. a. V. B. S. Sheetal, "Knowledge Graph-Based Thematic Similarity for Indian Legal Judgement Documents using Rhetorical Roles," India, 2018.

[3] J. S. L. O. a. T. S. M. Edge, "From Local to Global: A Graph RAG Approach to Query-Focused Summarization of Legal Norms," 2024.

[4] Y. X. X. G. K. J. J. P. Y. B. Y. D. J. S. M. W. a. H. W. Y. Gao, "Retrieval-augmented generation for large language models: A survey," China, 2023.

APPENDICES

Appendix 1. Work Breakdown Structure (WBS)

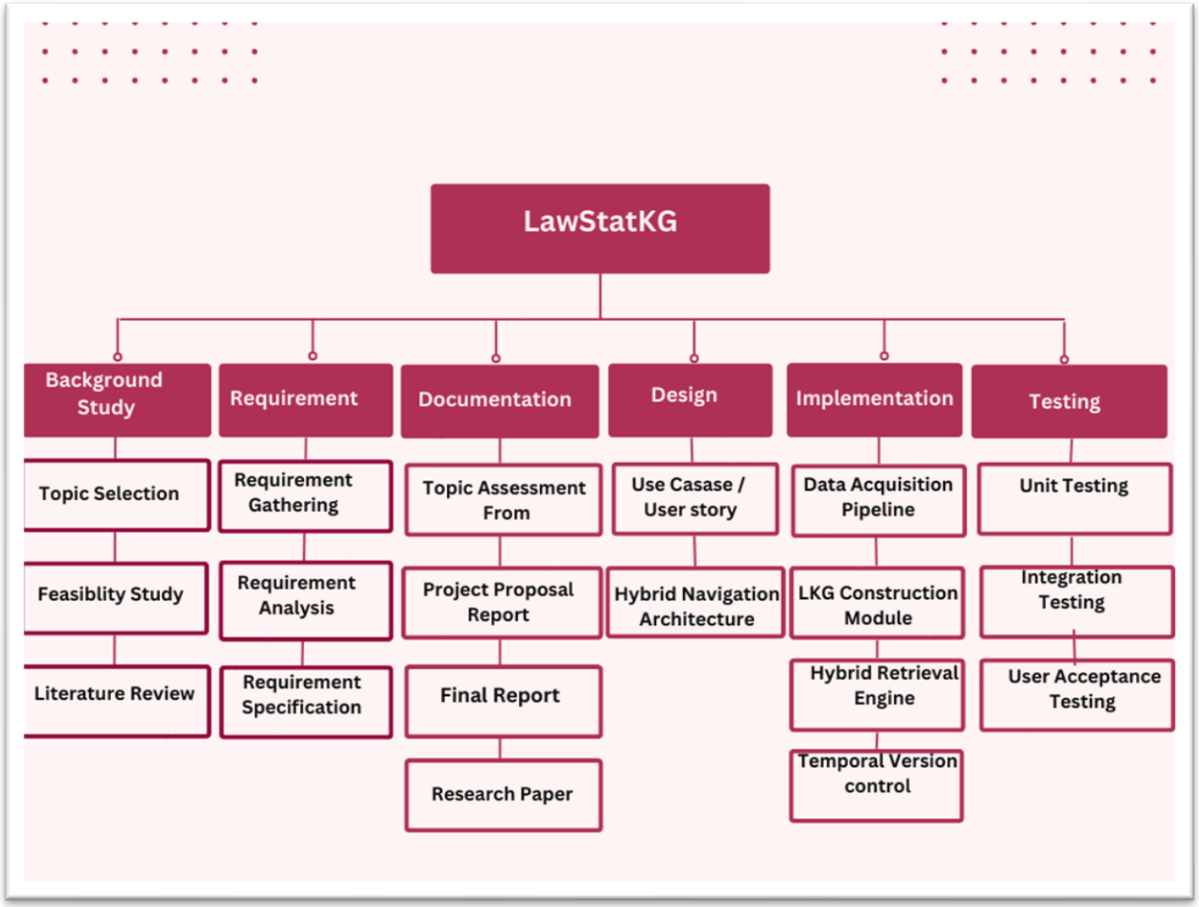

Appendix 2. Gantt Chart

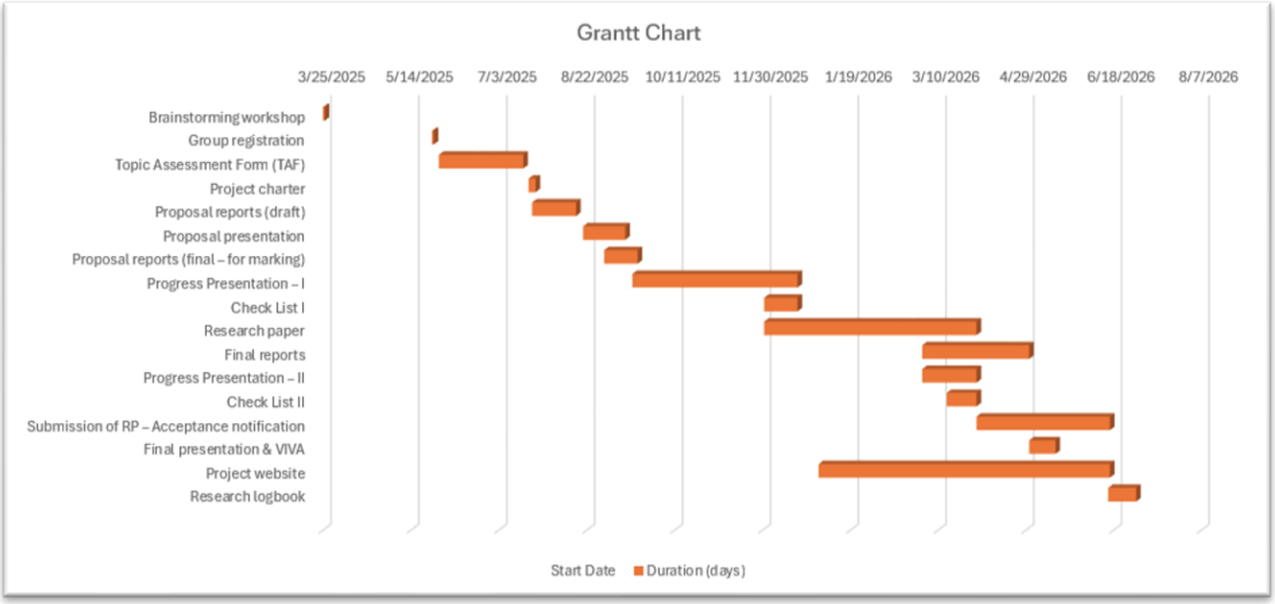

Appendix 3. Similarity Report

Proposal draft

ORIGINALITY REPORT

5%	4%	1%	4%
SIMILARITY INDEX	INTERNET SOURCES	PUBLICATIONS	STUDENT PAPERS

PRIMARY SOURCES

1	Submitted to Sri Lanka Institute of Information Technology Student Paper	2%
2	louisdl.louislibraries.org Internet Source	1%
3	Submitted to The Robert Gordon University Student Paper	<1%
4	Submitted to UCL Student Paper	<1%
5	Submitted to The Shah Noorani Institute of Technology (SNIT) Student Paper	<1%
6	www.tdx.cat Internet Source	<1%
7	www.coursehero.com Internet Source	<1%